

Mobile Device Security

- Design assumes:
 - Someone steals your iPhone
 - Device is passcode-protected and locked at time of theft
- Potential attacks:
 - Exhaustive search for passcode
 - Impersonate fingerprint / face
 - Take apart phones and remove flash storage / read from RAM
 - Exploit a bug in the OS kernel
 - USB / WiFi / radio might have bugs to exploit
 - Install a hacked version of the OS w/o security checks
 - I.e. redirect network traffic to use adversary's update server
 - Take apart phone and write different OS to flash chip

iOS Hardware Architecture

- Components
 - Main CPU
 - DRAM
 - Flash file storage
 - AES encryption engine between main CPU / DRAM and flash storage
 - Secure enclave CPU
 - UID cryptographic key in enclave hardware
 - We can use this to decrypt / encrypt but we cannot get this back out
 - Authentication sensors: fingerprint reader / face ID camera

Secure Boot

- Goal: make sure adversary cannot run modified OS / apps / enclave
- At power on, CPU starts executing from boot ROM that is read-only
- Boot sequence: boot ROM -> iBoot -> OS kernel -> apps
 - Boot ROM: checks iBoot code is signed by Apple's secret key
 - Checks it using Apple's public key
 - iBoot: checks OS kernel code is signed by Apple
 - OS kernel: checks apps are signed by developer
- How to check things are signed:
 - Apple holds some private key only known to them and code is signed
 - Code is hashed
 - This hash is then encrypted with Apple's private key
 - This hash is then decrypted using the public key stored in the secure enclave, then hash of the entire code has to match this decrypted hash

Downgrade Attacks

- Downgrade attacks:
 - Replace OS kernel with an older one with a known bug
 - Bad solution:
 - iBoot checks OS kernel version is at least a certain version
 - Adversary could modify the stored version number
 - Good solution:
 - Each device has a unique ID called the ECID
 - Apple update servers signs each upgrade with a specific ECID

- Boot sequence checks signature is for this phone's ECID
- Good because:
 - Apple update servers will not sign an out-of-date kernel
 - Adversary unlikely to have old software with victim's exact ECID

Secure Enclave

- Purpose: prevent main CPU from ever seeing crypto keys
 - Also defend against passcode guessing
 - CPU asks enclave to do things i.e. help decrypt certain things
 - Enclave is a separate CPU with a secure boot sequence
 - Shares DRAM with the main CPU
 - Encrypts its own memory contents
 - Runs fixed Apple software
-
- Overall plan: all user data is encrypted
 - Allow decryption only phone is unlocked
 - Offloaded to separate CPU to prevent compromised OS kernel from getting decryption keys / fingerprint / face data or bypassing passcode retry limits
 - Allows hiding of keys with memory encryption / auth so direct read of RAM doesn't reveal keys
 - The main CPU cannot have all memory encrypted, but secure enclave is fine
 - The decryption keys are computed based on passcode
 - More on this later, what ends up happening is only three or so keys are computed based on this and are used to wrap other keys
 - This is cached
 - Fingerprint / face recognition data can be used to retrieve cached value
 - We avoid having the OS communicate the sensor data directly to the secure enclave
 - Instead the sensor data is encrypted with manufacturing time secret key that the secure enclave knows

Data Encryption

- Files are encrypted in flash storage
 - They generate data encryption keys based on the passcode rather than storing it permanently
 - This way they exist nowhere on device after reboot and can be forgotten when device locks
 - The encryption key is based on both the UID (hardware) of the secure enclave and the password
 - The key is generated slowly to limit speed of exhaustive search, and enclave can also limit attempts
 - Dependent on UID to avoid brute forcing outside of the device
 - Background apps:
 - When the phone is locked, some apps have to be able to run and read/write files
 - OS kernel can generate a secret key for each of these apps and then "wrap" the key with it for just this file and give it
 - The original actual key is no longer actually stored and only this app should be able to figure out what the actual key they need is
-
- The entire filesystem metadata is encrypted with a key that is based on the UID so that even if you take this flash chip out and put it in a different phone, it cannot read
 - This is given out to the OS once the phone boots
 - Each separate file is encrypted with a different key, which is then wrapped by another key called the kdf before being sent back to the OS
 - Are stored in the secure enclave only
 - These keys can then be allocated out to different parts of the OS for different protection levels that are based on when the OS deletes the keys from its memory
 - Complete: can decrypt only when phone is currently unlocked

- Derived from passcode on unlock, discarded when locked
- Complete unless open: file can be written at any time but not read, used specifically for downloading attachments
- Until first authentication
 - When phone was unlocked for the first time, this is derived
 - Discarded only on power-off
 - Default for third-party app data
- No protection: can decrypt any time: only derived from UID and not passcode
- After the OS deletes the key, it will have to go back to the secured enclave with the password and re-ask for the key, which will require the password
- To use these wrapped file keys:
 - FS code on CPU asks secure enclave to get certain keys for certain files
 - CPU never receives raw keys, but instead wrapped ones
 - The secure enclave and the AES engine (between the CPU and its flash memory) share a secret
 - The CPU then has to send this to the AES engine to "program" it, so the CPU still never sees the actual key
- This allows you to change the pin fast:
 - You just need to recompute these kdf keys (only 3-4) and then you're good because the actual file keys don't have to change
 - Secure enclave will then just have to start encrypting the filesystem keys with new stuff