

Network Security

- A big open network invites many attacks
 - Authentication / confidentiality
 - Liveness (uptime of services)
 - Privacy
 - Today: attacks on core network protocols
-

- Internet core:
 - Open network consisting of many ISPs
 - Peering agreements between ISPs provide connectivity
 - Many protocols:
 - TCP, BGP, DHCP, UDP, telnet, FTP, IMAP
 - Defined by RFC interfaces
 - Many implementations of these
 - Changing the protocol is much harder than changing implementation
- There are now many secure higher layers such as:
 - Kerberos, SSH, SSL
 - Enabled by the advent of cryptography
- Core network still lacks security
 - Mostly concerned with liveness
 - Hosts manage end-to-end security instead of the actual network
 - This design is successful, but important to look at attacks

Forging IP Addresses

Remote Login

- Example internet application: remote login (1980)
- Uses telnet to open a TCP connection to the login program
- Attacker can:
 - Steal the password by snooping on network
 - Modify data in flight
 - Inject false data
 - Re-direct conversation with routing
- rlogin deals with this by not using a password to authenticate
 - Destination only allows hosts in its hostfile to log in
 - Attacker could pose as the client's IP address
 - But TCP communication is both ways, so if the attacker lies about the source, then the server's replies will not go back to the attacker

TCP Setup

- On initial handshake, the client sends a sync request to the server with its sequence numbers
- Server responds with its sequence numbers and an acknowledgement of the client sequence numbers
 - Ack is important so the client knows that this is coming from the real server
 - The client only has a guarantee that its destination packet reached the correct server
 - Does not have guarantee that its receiving messages from the real server
 - Client sends back an acknowledgement of the server's sequence numbers and the data it wants to send

TCP Sequence Number Attack

- Adversary wants to simulate a connection to the server by a client C
- Attacker sends a TCP packet with src = C
- Server will send sequence numbers to C
- Attacker has to guess the sequence numbers sent to C to be able to communicate with the server
 - Could try guessing based on guessing how the server does sequence number management
 - Attacker could do an actual connection to the server to learn about what kinds of sequence numbers the server is currently using and then guess the next one
- The real client will receive the actual packet and might think that it is an old packet
 - Real client might send an abort packet to shut down this connection
 - The attacker might have to race to get there before the RST arrives
- Security risks:
 - If authentication is based on IP address, SN attacker can pretend to be a host in the trusted user list
 - Breaks rlogin; key problem was assuming that the TCP layer's sender field was legit
 - IP authentication is no longer used for remote login
 - Denial of service attack:
 - Can constantly send reset packets
 - Servers actually accept reset packets for a large window of sequence numbers so you don't even need to get the sequence number exactly right
 - You can just guess
 - These can be used to target TCP connections between BGP routers
 - Causes routers to assume link failure and affect traffic
 - Hijack existing connections
 - Can inject data into an existing connection
 - Wait for someone to log in and then hijack
- Mitigation:
 - End to end encryption (i.e. SSL as used in the next lecture)
 - ISPs can filter packets with obviously forged IP source addresses
 - Hard for complex networks

Hardening TCP against forged IP source addresses

- Need to make it harder for attacker to guess next ISN (initial sequence number)
- Can't choose them completely randomly or else you violate TCP spec
 - Need to avoid recently used sequence numbers for same host / port pair
- Can't do random increments because we don't want to wrap around very often
 - Need to keep these increments small
- Have to be careful about how we generate random numbers since random number generators can be reverse-engineered
- Main idea that is implemented in most OS is to use a random offset for each src/dest pair
 - $ISN = ISN_old + SHA1(srcip, srcport, dstip, dstport, secret)$
- Whole idea is to prevent attacker from being able to make an ordinary connection and using that info to try to guess the ISN for another client

Liveness

SYN Attacks

- SYN flooding: first high-profile DoS attack
 - SYN = synchronization request (initial packet)
 - Server must check client's ACKs to it sending back sequence number
 - The original implementation of TCP kept the states for these "half-open" connections
 - Kept it for minutes if client is slow / network lossy
 - Only willing to remember 50 half-open connections to avoid OOM
 - Silently ignored new connections if it already had 50 waiting

- Attack sends SYN packet with forged random IP addresses
 - Most of these will never respond and the server just waits for them to send a third packet when they never will
 - Server begins to ignore legitimate connection requests
 - Don't even need to send these that fast because servers keep half-open connections for minutes
-

- Mitigation: SYN cookies
 - Make the server stateless until it receives the third packet
 - Tricky because the half open-state was what helped ensure source IP wasn't forged by checking the packet had the right ACK
 - Use a bit of cryptography
 - Encode server-side state into sequence number
 - $ISN = SN_c + \text{timestamp} \parallel \text{SHA1}(\text{src/dest addr/port, secret, timestamp})$
 - Timestamp is coarse-grained (minutes) to ensure that clients have enough time to submit back a packet
 - ISNs are per client so attacker can't guess for a forged IP address
 - Upon receiving the third packet, the server just has to recompute to see if this is a feasible packet which should have only been known by the client that receives it legitimately
 - Successfully blunted low-rate SYN-flooding DoS attacks

Bandwidth Amplification

- Attacker's goal is to overwhelm server / link so legit traffic is discarded
 - Send ICMP echo request (ping) packets to broadcast address of a network
 - Used to be that you would get an ICMP echo reply from all machines on network
 - If you fake a packet with a victim's address, then the victim gets all replies
 - If you find a subnet with 100 machines on a fast network, then you get 100x amplification on your attack
 - Fixed by routers blocking directed broadcast (packets sent to broadcast address)
-

- Modern variant:
 - Use DNS because with a small query, a server might send back a large responses
 - DNSSEC makes this even worse since it contains lots of signatures
 - DNS runs over UDP so source address is completely unverified
 - Hard to fix this because DNS has to respond to anyone
- Generally, can be mitigated via a DoS protection service (i.e. Akamai)
 - These have lots of available bandwidth and you tunnel requests into them first
 - They can filter out DoS attack traffic (i.e. DNS amplification responses)

Routing Protocol Attacks

- Generally, routing protocols are overly-trusting of participants
- ARP: within a single Ethernet network sending via MAC addresses of routers
 - To send an IP packet, you just need the MAC address
 - In this protocol, you broadcast a request for target's mac and anyone can listen to / send a reply with no authentication
 - You can just impersonate people
- DHCP: within a single Ethernet network sending via IP address
 - Client asks for IP address by sending a broadcast request
 - DHCP server responds with no auth
 - Adversary can impersonate DHCP server to new clients on the network
- BGP: Internet-wide
 - Huge routing system
 - Any BGP participant router can announce route to any IP address

- You can announce you have a path to MIT and people will route through you
 - You can inspect / modify traffic and then forward to MIT
- Fixes:
 - Trusted database of who is allowed to announce what IP prefixes
 - This database is a weak point
- The open Internet makes it easy for attackers to gather useful info
 - Can probe random hosts to see if they are running vulnerable software / protocols
 - Can scan entire Internet (2^{32} addresses) in just 45 minutes

Improving Security

- Firewalls: partial fix but widely used
 - Issue if adversary is within firewalled network
 - Hard to determine if packet is malicious
 - Hard to authenticate src / dst fields
- Cryptographic security on top of TCP / IP can be used but is hard