

- So far: isolation
- Upcoming: case studies about how systems have used isolation

OpenSSH

- Non privilege separation version:
 - On a host computer, sshd runs on port 22
 - Runs with root privileges
 - When it receives a tcp connection, it forks and creates a pseudoterminal with the user id of the corresponding user
 - The connection talks with this to do bash stuff
 - This also has to run with root permissions
 - Has to access the host key to sign messages and use user password
 - Does all of the compression / decoding / cryptography
-

- Principle of Least Privilege
 - Every component has least privileges to do its job
 - Challenges:
 - How to split / isolate / share?
 - How to maintain performance?
 - How hard is it to change the code?
-

- OpenSSH privsep:
 - sshd still runs as root, but when it forks it creates a per connection monitor
 - This spawns a worker that actually maintains the TCP connection
 - After the connection is made and authentication is made, another worker is spawned which only has user permissions and handles making the pseudoterminal
 - Slight technicality:
 - Three process plan: the TCP connection worker stays around and just proxies to the other work
 - Two process plan: kill the TCP connection worker and reroute the TCP connection to the worker
 - OpenSSH actually does this b/c it has performance improvement
 - The monitor process provides certain things such as encrypting messages / authenticating user
 - The root monitor process is now no longer exposed directly to the user
 - The monitor also has an FSM that limits the interface provided to the user process to further increase security
 - Security of various components:
 - sshd listener: root
 - monitor: root
 - pre-auth worker: anonymous UID, has access to sign message as host computer and send stuff to computers on the network
 - However, you can't use this for new connections because those would require using a new host key
 - You can only use this for existing connections that used that same host key
 - post-auth worker: could act as user, do DOS attack
 - Attack surfaces:
 - sshd listener: creating new TCP connections, not even reading the TCP connections
 - monitor: interface to worker
 - pre-auth worker: reading / writing to TCP connections (most likely to be exploited)
 - Not leaking any user data / root access
 - post-auth worker: pre-auth worker state / TCP data, but this only occurs if you can log in
 - All of the bugs found in the past fell in the pre-auth / post-auth worker
-

- Challenge-response protocol:

- A process where the server sends a challenge c to a client and the client has to sign it with their secret key to prove their identification
 - The server then uses the client's public key to verify that this is in fact the client they expect
 - Client's secret key: stored in `~/.ssh/id_rsa`
 - Server's authorized public key: `~/.ssh/authorized_key`
 - You cannot have the worker do this process because we assume the worker could be bad
 - The monitor has to generate challenges and send it to the client
 - Otherwise the worker could pretend that it satisfied the challenge or use a replay attack with an old challenge
-

- Spawning post-auth worker
 - Need to send all of the low-level state to the post-auth worker to make the client seem like they are talking to the same process
 - Most tricky part of the implementation