

Private Information Retrieval

Linear homomorphism turns out to be a useful property for building cryptosystems that can do operations on encrypted data

Today, we will be looking at querying remote databases without revealing what we are searching for

- Suppose a server holds an N bit database $D \in \mathbb{Z}_2^N$
- A user has an index i and wants to discover D_i without revealing i to the server

Private Information Retrieval Scheme

A PIR scheme is a triple of randomized algorithms (`Query`, `Answer`, `Reconstruct`) that meet the following properties, where N is the length of the database:

- **Correctness:** for any index i , database D , and security parameter n , the probability of `Reconstruct(Answer(D, qu), st)` is at least $1 - \text{negl}(n)$ where `(qu, st)` are given by `Query(n, i)`
- **Security:** for any two indices i, j , the queries corresponding to i and j are computationally indistinguishable
- **Succintness:** The total bit length of the client's query and the server's answer is less than N
 - That is, the total number of bits communicated is smaller than N
 - This is to rule out trivial schemes where the user downloads the whole database

Square-Root PIR

Key idea is that the server will represent its N bit database as a matrix $D \in \mathbb{Z}_2^{\sqrt{N} \times \sqrt{N}}$

- The user wants to read bits at 2D locations (i, j)
- Assume we are given a linearly-homomorphic encryption scheme (`Gen`, `Enc`, `Dec`)

Protocol:

- `Query($n, (i, j)$)`:
 - Build the unit vector u_j that consists of all 0s except for a single 1 at position j
 - Sample a secret sk with `Gen(n)`
 - Output the query vector $\text{qu} \leftarrow \text{Enc}(\text{sk}, u_j)$ along with client-held state $\text{st} \leftarrow (\text{sk}, i)$
- `Answer(D, qu)`:
 - Output the answer vector $\text{ans} \leftarrow D \cdot \text{qu}$
- `Reconstruct(ans, st)`:
 - Decrypt the answer vector as $v \leftarrow \text{Dec}(\text{sk}, \text{ans})$
 - Parse the client-held state (sk, i)
 - Output the i th entry of v

Note that `Query` will run on the client, not the server

- The server is only in charge of addressing `Answer`

Correctness:

- If the encryption scheme is linearly homomorphic, we have that `ans` is:
 - $D \cdot \text{Query}(n, (i, j)) = D \cdot \text{enc}(\text{sk}, u_j) = \text{Enc}(\text{sk}, D \cdot u_j)$
 - Decoding and getting the i th bit will then get D_{ij}

Security:

- Follows from the CPA-security of the encryption scheme

Succintness:

- `Query` consists of \sqrt{N} ciphertexts, as does the output of `Answer`

Naive Implementation

We attempt to implement with the secret key encryption scheme from LWE

- **Query :**
 - We build the unit vector u_j as described above
 - We sample:
 - A random matrix $A \in \mathbb{Z}_q^{\sqrt{N} \times n}$
 - A secret key vector $s \in \mathbb{Z}_q^n$
 - A random error vector $e \in \chi^{\sqrt{N}}$
 - Output:
 - Query: $(A, A \cdot a + e + \lfloor q/2 \rfloor \cdot u_j)$
 - Client-held state: (s, i)
- **Answer :**
 - Parse the query as (A, b)
 - Output the answer $(D \cdot A, D \cdot b)$
- **Reconstruct :**
 - Parse the client-held state as (s, i)
 - Parse the answer as (H, c)
 - Compute $v = c - H \cdot s$
 - Round the i th entry of v to the nearest multiple of $\lfloor q/2 \rfloor$ and divide by $\lfloor q/2 \rfloor$

In practice we use:

- $n \approx 1024$
- $q \approx 2^{16}$
- Total communication cost would be $10^4 \cdot \sqrt{N}$ bits

Optimizations

- We can use the same A matrix as long as we use a random s and e each time
 - We can even use A as a public parameter known to all users
 - The A matrix now no longer needs to be sent around
- The server can precompute $D \cdot A$
- The user can prefetch the value of $D \cdot A$

These make the communication size independent of n

- The PIR scheme now uses $\log q \cdot \sqrt{N}$ bits, which is around $16\sqrt{N}$ bits