

NonBayesian Hypothesis Testing

Previously, we looked at a Bayesian approach where we had a *a priori* probability distribution and cost criterion to determine η

- We computed the likelihood ratio $\frac{p_{y|\mathcal{H}}(y|H_1)}{p_{y|\mathcal{H}}(y|H_0)}$
- However, it is not quite clear how to choose this
- The threshold η can be quite sensitive to these choices
- We'll go over another way to decide between hypotheses

Classification Performance Measures

The performance of any decision rule can be fully specified by two quantities:

$$P_F = P(\hat{H} = H_1 | H = H_0)$$

$$P_D = P(\hat{H} = H_1 | H = H_1)$$

- P_F is the false-alarm probability
- P_D is the detection probability

These have other names in other disciplines

- Statistics: P_F is the size and P_D is the power
- Type I error: P_F
- Type II error: $1 - P_D$
- Machine learning: P_D is the recall / sensitivity
 - Use $P_P = P(H = H_1 | \hat{H} = H_1)$ as the precision / positive predictive value (PPV)
 - This does depend on the choice of priors, but given a fixed set of these priors, it is a function of P_F

Operating Characteristic of the Likelihood Ratio Test

In general, a good decision rule has a large P_D and a small P_F

- However, these are competing objectives
- If we use a likelihood ratio test, then each choice of η uniquely determines a decision rule with some $(P_F(\eta), P_D(\eta))$
- We can trace this point as η varies in $[0, \infty)$ to get the **OC-LRT**, also known as the **ROC (receiver operating characteristic)**

In general, we have that:

- As $\eta \rightarrow 0$, we approach $(1, 1)$
- As $\eta \rightarrow \infty$, we approach $(0, 0)$

Bayesian Operating Point

The Bayesian test corresponds to a specific η , which means it is a point on the curve

To get this, we note the Bayes risk is of the form:

$$\phi(\hat{H}) = \sum_{i,j} C_{ij} P(\hat{H} = H_i | H = H_j) P_j$$

This corresponds to:

$$\phi(\hat{H}) = \alpha P_F - \beta P_D + \gamma$$

$$\phi(\hat{H}) = (C_{10} - C_{00})P_0 \cdot P_F - (C_{01} - C_{11})P_1 \cdot P_D + C_{00}P_0 + C_{01}P_1$$

For any constant c , setting $\phi(f) = c$ yields a line:

$$P_D = \frac{\alpha}{\beta} P_F + \gamma - c$$

In this case, $\frac{\alpha}{\beta}$ is equal to the threshold η_B in the Bayesian test

- Therefore, the Bayesian test is finding the smallest c that leads to an intersection of the line with slope η_B and the OC-LRT
- Since the y-intercept contains a $-c$, this means that smaller c raise the line upwards
- This means we are looking for the highest line with slope η_B that still intersects the OC-LRT

Monotonically Nondecreasing

The OC-LRT is monotonically nondecreasing

- This can be proven by considering that if you increase η such that P_F changes, then P_D and P_F cannot increase by looking at the definitions

Neyman-Pearson Hypothesis Testing

When we are unable / it is difficult to choose costs for a decision problem, one popular alternative is:

- Choose a decision rule that maximizes P_D
- Subject to $P_F \leq \alpha$ for some α of choice

Neyman-Pearson Lemma

Suppose the likelihood ratio $L(y)$ is a purely continuous random variable under each hypothesis. Then a solution to the above among the deterministic estimates is an LRT test with η chosen such that:

$$P_F = \alpha$$

What this implies is we can just look at the OC-LRT and draw the line P_F *alpha* to find the P_D value

This theorem is quite restricting in that it:

- Requires the likelihood ratio is continuous, which is not true if either hypothesis uses a discrete distribution or if there is a strange interaction
- Only looks at deterministic decision rules

Proof Sketch:

- We look at any test with P_D, P_F as well as a specific test under a LRT with threshold η with P_D^η, P_F^η
- We can do some algebra to find that:

$$P_D^\eta - P_D \geq P_F^\eta - P_F$$

- As a result, if η is chosen such that $P_F \leq P_F^\eta$, then $P_D^\eta \geq P_D$
- If η is chosen such that $P_F^\eta = \alpha$, then any other test with $P_F \leq \alpha$ will have $P_D \leq P_D^\eta$, and therefore is not better

P-Value

The LRT rule $L(y) \geq \eta$ can have a monotonic $g(\cdot)$ applied to both sides:

$$g(L(y)) \geq g(\eta)$$

We can do this with P_F , which gives:

$$P_F(L(y)) \leq P_F(\eta)$$

- We take $p_*(y) = P_F(L(y))$ and $P_F(\eta) = \alpha$

- This is known as the p-value and the significance level of the test