

Causal Inference

A linear relationship between X and Y does not imply causation

Causation implies:

- Changing X necessarily changes Y
- Typically a time component, so the change in X has to occur before the change in Y

Counterfactual Model

Let $X \in \{0, 1\}$ be a random variable denoting whether or not a treatment was applied

- $X = 0$ is the "control"
- $Y \in \mathbb{R}$ is the response, which is typically a binary response but really could be any real

We define C_0 and C_1 to be the potential outcomes, i.e. the potential values of Y in the case that $X = 0$ or $X = 1$

- These are random variables
- $Y = C_X$
- When $X = 0$, C_0 is the observed, and C_1 is the counterfactual. Vice versa for when $X = 1$

☰ Average Treatment Effect

$$\theta = \mathbb{E}[C_1] - \mathbb{E}[C_0]$$

- Generally, we do want to find the ATE, but it is not computable directly because we don't know the counterfactual for each observation

☰ Association

$$\alpha = \mathbb{E}[Y|X = 1] - \mathbb{E}[Y|X = 0]$$

- We can estimate this directly through the plugin estimator

In general, these are not equal to each other due to the idea of confounding variables

- Consider the example where we have $Z \sim \text{Unif}([-1, 1])$ and let $X = 1(Z > 0)$
- We let $C_0 = C_1 = Z$
 - Clearly, the treatment has no effect, and we get that $\theta = \mathbb{E}[C_1] - \mathbb{E}[C_0] = 0$
 - However, $\alpha = \mathbb{E}[Y|X = 1] - \mathbb{E}[Y|X = 0] = 1/2 - -1/2 = 1$
- Z is a confounding variable that impacts both X and Y
 - An example of another confounding variable is the decision to give someone a treatment or not based on their current health

Estimating ATE with Association

- We can design an experiment to guarantee $\theta = \alpha$
 - Known as randomized controlled trials, or A/B testing
 - This is done by determining who gets which treatment entirely by a random coin
 - This can be unethical, i.e. denying someone with a medical condition a treatment based on a flipped coin
- As long as X is a pure Bernoulli random variable, then we guarantee that $\alpha = \theta$
 - This is because C_0 and C_1 are independent of X

-
- Often, we use observational studies where the outcomes are not independent of X

- However, we might be able to find a feature vector Z such that C_0, C_1 are independent of X conditioned on Z
 - I.e. in a treatment, we could have Z be a vector of all of your health conditions
- We could then split up our observations based on the value of Z , and for each of these, we would have that $\theta = \alpha$
- However, this is hard to do when Z is high-dimensional
 - You would have to create bins for each Z , but then each bin would only have a small number of observations
 - We want to somehow reduce Z to a single number, which is much easier to bin up

Propensity Score

$$p(z) = P(X = 1|Z = z)$$

- If we calculate this for each Z value and then bin them up then we can apply the following theorem:

Rosenbaum and Rubin

If C_i is strongly unconfound with X given Z , then it is also strongly unconfound with X given $p(Z)$

Typically, we can't compute $p(z)$ directly, but we could learn it via logistic regression